

Real-Time Location Tracking as an AI Monitoring Application: A Comparison of Rideshare and Dealership Shuttle Services

Course: ARTI 409: AI Management and Maintenance

Instructor: Tim Cruz

Submitted by: Roy Aggarwal

Date: July 12, 2026

Introduction

Real-time location tracking has become a standard feature of consumer transportation services, but the technology behind the map is not always the same. This case study compares two real-time tracking experiences: the Uber rideshare app and a dealership's courtesy shuttle tracking link. Both display a moving vehicle on a map, but the underlying systems differ in a way that affects reliability and user trust. This distinction matters because it illustrates a broader principle in AI monitoring and optimization: raw sensor data, such as GPS coordinates, becomes far more useful once a predictive model is layered on top of it to correct for noise, delay, and real-world conditions.

The significance of this comparison extends well beyond personal convenience. Ride-hailing platforms operate at a scale where they must process billion-level volumes of location and arrival-time requests every week [2], and the global user base for these platforms is projected to reach 2.34 billion people by 2030, with a user penetration rate of roughly 24.6 percent in 2026 [4]. At that scale, even small gains in prediction accuracy translate into a meaningful difference in service quality across an enormous and still-growing population of riders. Any organization that operates a vehicle fleet, whether a global rideshare platform or a small dealership shuttle, depends on accurate and timely location data to manage customer expectations, build trust, and support day-to-day operations. Understanding what separates an AI-enhanced tracking system from a basic one provides insight into where smaller operators could realistically improve their own services without needing platform-scale resources.

AI Technologies Analysis

Uber's real-time tracking is built on a machine learning system rather than a raw GPS feed alone. A conventional routing engine calculates an estimated arrival time using a shortest-path algorithm applied to a digital road network, but this approach cannot fully account for real-world conditions such as traffic, weather, or driver behaviour. Uber addresses this by layering a deep learning model, trained on historical trip data and real-time signals, on top of the routing engine's initial estimate, refining it toward the actual outcome [1], [2]. This model, a residual neural network architecture, was built

specifically to replace an earlier production system based on XGBoost, a decision-tree ensemble method that had been the prior standard for this type of tabular prediction task [2]. The model is trained on a dataset of roughly 1.4 billion ETA requests, incorporating contextual, temporal, geospatial, trip-level, and traffic features, and applies a specialized loss function designed to stay robust when individual predictions are unusually far off, since large errors have an outsized negative effect on the rider experience [2]. This model functions as a post-processing step applied at scale across the platform's global network [2]. In practice, this means the moving vehicle shown on screen is not a direct readout of a GPS coordinate; it reflects a continuously corrected prediction, which is why the map updates smoothly without any action from the rider.

The dealership shuttle tracker appears to work differently, based on my direct observation of the service. The following is an assumption, since the dealership did not disclose its technology stack. The tracker was a link sent by text message that displayed the shuttle's position accurately, but the position only updated after I manually refreshed the page. This behaviour is consistent with a basic GPS tracking architecture, where a GPS module reports coordinates that are displayed on request rather than pushed continuously to the client [3]. There is no indication of a predictive or correction layer sitting between the raw GPS feed and the screen, and no equivalent of a trained model reconciling the reported position against expected conditions. This is a reasonable assumption given the observed behaviour, but it should be treated as an inference rather than a confirmed fact about the dealership's actual system.

The core technical distinction, then, is not the presence of GPS hardware, which both systems likely have, but whether a machine learning model sits on top of that raw data to predict, smooth, and continuously push an updated position. Uber's system does this, using a purpose-built architecture trained on a massive, feature-rich dataset. The dealership's shuttle tracker, as observed, appears to rely on the raw signal alone.

Impact and Efficiency

The practical impact of this difference is measurable in both accuracy and operational efficiency. Uber's ML-based post-processing has been shown to meaningfully outperform naive, non-ML ETA estimates in both offline and live production testing [2]. This performance is achieved while still meeting the demands of real-time service: reported latency figures show a median response time of roughly 3.25 milliseconds and a 95th-percentile response time of about 4 milliseconds even while handling billion-level weekly request volumes [2]. This combination of improved accuracy and consistently low latency at massive scale is what allows the rider-facing map to feel both correct and instantaneous.

This accuracy gain directly affects rider trust: an ETA that is consistently close to the actual arrival time reduces uncertainty and allows riders to plan the last few minutes before a pickup with confidence. From my own experience, the difference was noticeable in a practical sense rather than a technical one. With Uber, I could glance at

my phone at any point and know exactly where the driver was without doing anything. With the dealership shuttle, I had to remember to manually refresh the page to get an updated position, which meant the information I was looking at could already be several minutes stale. For a rider trying to time when to leave their house or step outside, that gap has a real effect on convenience, even though both systems technically provided the correct location when checked.

Challenges and Solutions

A key challenge for any real-time tracking system is maintaining accuracy when GPS signals are weak or interrupted, such as in a parking garage or an area with poor cellular coverage. Uber's approach addresses this indirectly: because the ML layer is trained to predict a plausible position and ETA rather than only relaying a raw coordinate, it can smooth over brief GPS gaps and avoid displaying an erratic or frozen marker. A purely GPS-ping-based system has no such correction mechanism [3], so any gap in signal or any delay in the app polling for new data shows up directly as a stale or inaccurate display.

A second, closely related challenge is the risk of extreme, outlier errors that damage user trust more than typical small errors do. Uber's model design responds to this directly by using a loss function during training that is specifically built to be robust to outliers, rather than optimizing only for average-case accuracy [2]. This is a good example of a solution being designed around a known failure mode rather than treated as an afterthought.

A third challenge is one of infrastructure cost. Building and maintaining a machine learning pipeline of Uber's scale requires significant engineering investment, data volume, and ongoing model retraining, which is a realistic barrier for a small operator such as a dealership. A practical, lower-cost solution would not require replicating Uber's full system, but could still close much of the experience gap. For example, switching from a manually refreshed page to a push-based update model, where the app or webpage automatically polls for a new position every few seconds, would remove the need for the rider to take any action, even without adding predictive modelling. This alone would resolve the most noticeable pain point observed in this case study.

Conclusions and Implications

This comparison shows that the value of AI in a monitoring application is not always visible as a separate feature. In Uber's case, the machine learning layer is invisible to the rider; its effect is felt only as a smoother, more reliable experience, delivered at a scale and speed that would be difficult to achieve with rule-based methods alone. This suggests that AI-driven monitoring succeeds when it improves a system quietly, without requiring the end user to understand or interact with the underlying model.

For smaller operators such as dealership shuttle services, the implication is that meaningful improvements do not require adopting Uber's full machine learning infrastructure. Even a modest step, such as automatic client-side refreshing, would close much of the perceived experience gap at a fraction of the cost. Looking forward, as lightweight and pre-trained machine learning tools become cheaper and easier to deploy, it is reasonable to expect that mid-sized fleet operators, such as regional shuttle or delivery services, will increasingly adopt simplified versions of predictive tracking rather than continuing to rely on raw GPS pinging. Operators that do not adapt risk appearing outdated compared to AI-enhanced competitors, even where their underlying location data is just as accurate.

References (IEEE format compliant)

- [1] T. Binaykiya, X. Hu, and O. Cirit, "DeepETA: How Uber Predicts Arrival Times Using Deep Learning," Uber Engineering Blog, Feb. 10, 2022. [Online]. Available: <https://www.uber.com/en/us/blog/deepeta-how-uber-predicts-arrival-times/>
- [2] X. Hu, T. Binaykiya, E. Frank, and O. Cirit, "DeepETA: An ETA Post-processing System at Scale," arXiv:2206.02127, Jun. 2022.
- [3] A. H. Alquhali, M. Roslee, M. Y. Alias, and K. S. Mohamed, "IoT Based Real-Time Vehicle Tracking System," in Proc. 2019 IEEE Conf. on Sustainable Utilization and Development in Engineering and Technologies (CSUDET), Penang, Malaysia, 2019, pp. 265-270, doi: 10.1109/CSUDET47057.2019.9214633.
- [4] Statista, "Ride-hailing - Worldwide," Statista Market Forecast, 2026. [Online]. Available: <https://www.statista.com/outlook/mmo/shared-mobility/ride-hailing/worldwide/>